

B+ Trees

CSE 332 Summer 2021

Instructor: Kristofer Wong

Teaching Assistants:

Alena Dickmann	Arya GJ	Finn Johnson
Joon Chong	Kimi Locke	Peyton Rapo
Rahul Misal	Winston Jodjana	

Announcements

- ❖ Thank you SO MUCH for your patience this past week

- ❖ Exercise 6 out today: VerifyAVL
 - Due Sunday, 11:59 PM (released late)
 - Ex 7 & 8 will come out on time this weekend, they are on the easier side.

- ❖ Midterm coming next week!
 - Reminder: Non-traditional midterm

Lecture Outline

❖ Recap

❖ B+ Trees

- Goals and Design
- B+ Tree Structure
- B+ Tree Find
- B+ Tree Add
- B+ Tree Remove

Recap of weekend Lecture

- ❖ Our data structures so far have assumed $O(1)$ time for basic operations, reads and writes
- ❖ When our data structures are big enough, reads and writes may trigger a disk load (takes a LONG time)
- ❖ To mitigate this, we rely on locality
 - Spatial Locality
 - Temporal Locality
- ❖ We want a data structure that is specifically designed to take full advantage of locality and minimize disk accesses

Lecture Outline

- ❖ Recap
- ❖ B+ Trees
 - **Goals and Design**
 - B+ Tree Structure
 - B+ Tree Find
 - B+ Tree Add
 - B+ Tree Remove

Goal of the B+ Tree

- ❖ **Problem:** A dictionary with so many items *most of it is on disk*
- ❖ **Goal:** A balanced tree (logarithmic height) that minimizes disk accesses
- ❖ **Concept:** Increase the branching factor of our tree
 - Minimize number of nodes to traverse
- ❖ **Disclaimer:** You will not have to implement this structure!!
 - Requires more control over memory than Java allows


How do we minimize disk accesses?

- ❖ Increase size of each node in our tree
 - ... to the size of a full disk block
 - For a dictionary, this would mean many key/value pairs and pointers per node
- ❖ Worst case number of nodes that we look at for a find in any tree will be bounded by height
 - So let's try to maximize how efficiently we use the height
 - Higher branching factor than 2


Increasing efficiency: M-ary Search Tree

- ❖ A search tree with branching factor M (instead of 2)

- Each node has a key-sorted array of M children: `Node []`



- Ordering property similar to BST



- M-1 keys define the M subtrees (ie, ranges) that we search through
- ❖ Choose M such that the node size = disk block size

M-ary Performance?

❖ Runtime for `find` = `NumHops` * `WorkPerHop`

■ **Balanced** tree height is: $\log_M n$ (M-ary) vs $\log_2 n$ (binary)


- Eg: $M = 256$ ($=2^8$) and $n = 2^{40}$, M-ary makes 5 hops vs binary makes 40 hops

■ For each internal node, how to decide which child to take?

- Binary: Less than vs greater than node's single key? 1 comparison
- M-ary: In range 1? In range 2? In range 3?... In range M?
 - Linear search the `Node[]`: M comparisons
 - Binary search the `Node[]`: $\log_2 n$ comparisons

❖ Runtime for M-ary `find`:


■ $O(\log_2 M \log_M n)$



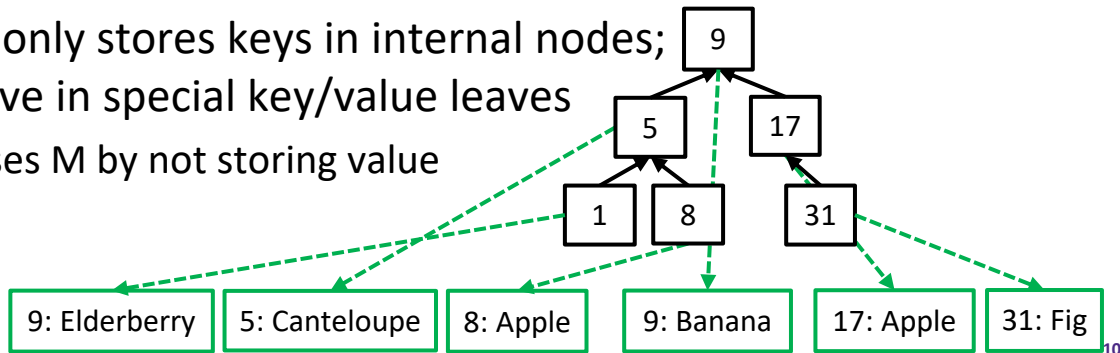
Design Decision: Key-only Internal Nodes

- ❖ A Dictionary ADT stores key->value pairs; where should we store a key's value?

- ❖ BST stores value alongside the key at every node
 - Loads entire node even if we are “passing through” to find a different key



- ❖ B+ Tree only stores keys in internal nodes; values live in special key/value leaves
 - Increases M by not storing value




Lecture Outline

- ❖ Recap
- ❖ B+ Trees
 - Goals and Design
 - **B+ Tree Structure**
 - B+ Tree Find
 - B+ Tree Add
 - B+ Tree Remove

B+ Tree Node Structure

Both the textbook and we refer to “B+ Trees” as “B-Trees”, but “B-Trees” actually encompass several variants

- ❖ Two node types: **internal** and **leaf**
- ❖ Each **internal node** contains up to $M-1$ keys (for up to M children)
 - Does not store values, only keys
 - Function as “signposts”
- ❖ Each **leaf node** contains up to L items
 - Stores (key, value) pairs
 - As usual, we’ll ignore the “along for the ride” value in our examples



3	“cat”
7	“apple”
21	“purple”
84	“ideas”

B+ Tree Parameters

- ❖ Two parameters, one for each type of node:

- M = # of children in an **internal** node
 - The ranges are defined by $M-1$ keys
- L = # of items in a **leaf** node

k_1	k_2	\dots	k_{m-1}	
ptr_1	ptr_2	\dots	ptr_{m-1}	ptr_m

(sorted by key)

- ❖ Picking M and L based on disk-block size maximizes B+ Tree's efficiency

- Recommend $M^* \approx \text{diskBlockSize} / \text{keySize}$
- Recommend $L = \text{diskBlockSize} / (\text{keySize} + \text{valueSize})$
- In practice, $M \gg L$
 - Since typically $\text{sizeof}(\text{key}) \gg \text{sizeof}(\text{keyvaluepair})$

k_1	v_1
k_2	v_2
...	...
k_L	v_L

(sorted by key)

* More precisely, we recommend

$$M = (\text{diskBlockSize} + \text{keySize}) / (\text{keySize} + \text{pointerSize})$$

B+ Tree Structure

❖ Internal nodes

- Have between $\lceil M/2 \rceil$ and M children; i.e., at least half full
- *Reminder: no values, just keys*

❖ Leaf nodes

- All leaves at the same depth
- Have between $\lceil L/2 \rceil$ and L items; i.e., at least half full
- *Reminder: keys **and** values*


❖ Root node – A Special Case!

- If tree has $\leq L$ items, root is a **leaf node**
 - Unusual; only occurs when starting up
- Else, root is an **internal node** and has between 2 and M children
 - i.e., the “at least half full” condition does not apply

B+ Trees are Balanced (Enough)

- ❖ Not hard to show height h is logarithmic in number of items n
 - Let $M > 2$ (if $M = 2$, then a “linked list tree” is legal – no good!)
 - Because all nodes are at least half full (*except possibly the root*) and all leaves are at the same level, the minimum number of items n for a height $h > 0$ tree is

$$n \geq 2 \lceil M/2 \rceil^{h-1} \lceil L/2 \rceil$$


minimum number of **leaves** minimum items per **leaf**

B+ Trees are Shallower than AVL Trees

- ❖ Suppose we have 100,000,000 items
- ❖ Maximum height of AVL tree?
 - Recall $S(h) = 1 + S(h-1) + S(h-2)$
 - $h = \mathbf{37}$
- ❖ Maximum height of B+ Tree with $M=128$ and $L=64$?
 - Recall $n \geq 2 \lceil M/2 \rceil^{h-1} \lceil L/2 \rceil$
 - $h = \mathbf{5}$

B+ Trees are Disk Friendly (1 of 2)

- ❖ Reduces number of disk accesses during `find`
 - Large M = shallower tree = potentially fewer accesses
 - Requires that we pick M wisely
 - Too large: multiple disk accesses to load a single **internal** node
 - Too small: tree could've been shallower
 - Time for binary search over $M-1$ keys insignificant compared to disk access
- ❖ Reduces unnecessary data transferred from disk
 - `find` wants one value; doesn't load "incorrect" values into memory
 - Only one disk access to bring (the single correct) value into memory: when we find the correct **leaf node**

B+ Trees are Disk Friendly (2 of 2)

- ❖ Maximizes temporal locality
 - B+ Tree-style **internal** nodes are used more often (they differentiate between a larger fraction of keys) than BST-style nodes, and therefore are more likely to be held in memory by the OS


Lecture Outline

- ❖ Recap

- ❖ B+ Trees
 - Goals and Design
 - B+ Tree Structure
 - **B+ Tree Find**
 - B+ Tree Add
 - B+ Tree Remove

B+ Tree Find/Contains

- ❖ M-way extension of a BST's root-to-leaf recursive algorithm
 - At each **internal** node, do binary search on (up to) $M-1$ keys to determine which branch to take
 - At the **leaf** node, do binary search on the (up to) L items
 - *Requires that keys are sorted in both **internal** and **leaf** nodes!*
- ❖ Difference:
 - Since we don't store value at internal nodes, we will never find our value in the internal nodes; must always traverse to the bottom of B+ Tree




Find/Contains Example

Notation:

- Internal nodes drawn horizontally
- Leaf nodes drawn vertically
- All nodes include empty cells

- ❖ Tree with $M=4$ (max # pointers in **internal node**) and $L=5$ (max # items in **leaf node**)
 - All **internal nodes** must have ≥ 2 children
 - All **leaf nodes** must have ≥ 3 items (but we are only showing keys)



Lecture Outline

❖ Recap

❖ B+ Trees

- Goals and Design
- B+ Tree Structure
- B+ Tree Find
- **B+ Tree Add**
- B+ Tree Remove

B+ Tree Add Algorithm (1 of 3)

1. Add the value to its **leaf** in key-sorted order
2. If the **leaf** now has $L+1$ items, *overflow*:
 - Split the **leaf** into two leaves:
 - Original **leaf** with $\lceil (L+1) / 2 \rceil$ smaller items
 - New **leaf** with $\lfloor (L+1) / 2 \rfloor = \lceil L/2 \rceil$ larger items
 - Attach the new **leaf** to its parent
 - Add a new key (smallest key in new leaf) to parent in sorted order

If step (2) caused the parent to have $M+1$ children, ...


B+ Tree Add Algorithm (2 of 3)

3. If step (2) caused an **internal node** to have $M+1$ children
 - Split the **internal node** into two nodes
 - Original **node** with $\lceil (M+1)/2 \rceil$ smaller keys
 - New **node** with $\lfloor (M+1)/2 \rfloor = \lceil M/2 \rceil$ larger keys
 - Attach the new **internal node** to its parent
 - Move the median key (smallest key in new node) to parent in sorted order
 - If step (3) caused the parent to have $M+1$ children, repeat step (3) on the parent
4. If step (3) caused the **root** to have $M+1$ children
 - Split the old root into two **internal nodes**, then add them to a newly-created **root** as described in step (3)
 - *This is the only case that increases the tree height!*

Add Example:

- ❖ Add 3, 18, 14, 30, 32, 36, 15, 16, 12, 40, 45, 38
- ❖ $M=3$, $L=3$


Add Example: Answer (1 of 7)



Special case: the
root is a **leaf node**

Add 3, 18, 14, 30, 32, 36, 15, 16, 12, 40, 45, 38
M=3, L=3

Add Example: Answer (2 of 7)

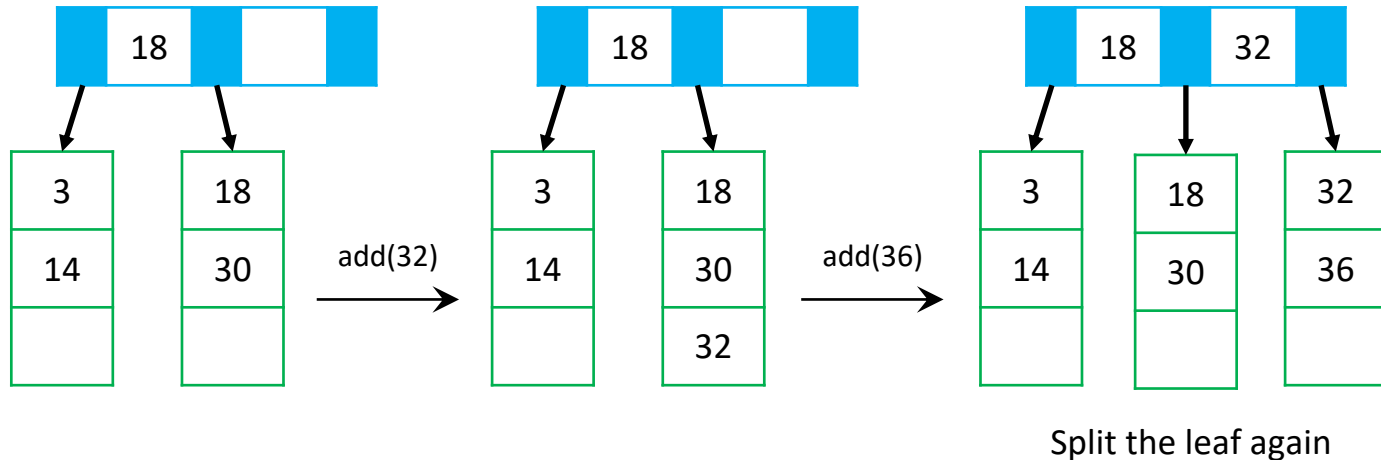


Special case: the
root is a **leaf node**

- When we “overflow” a leaf, it is split and the parent gains another key (to select between the two leaves)
- Parent’s new key is the smallest element in the right child
- If there is no parent, create one

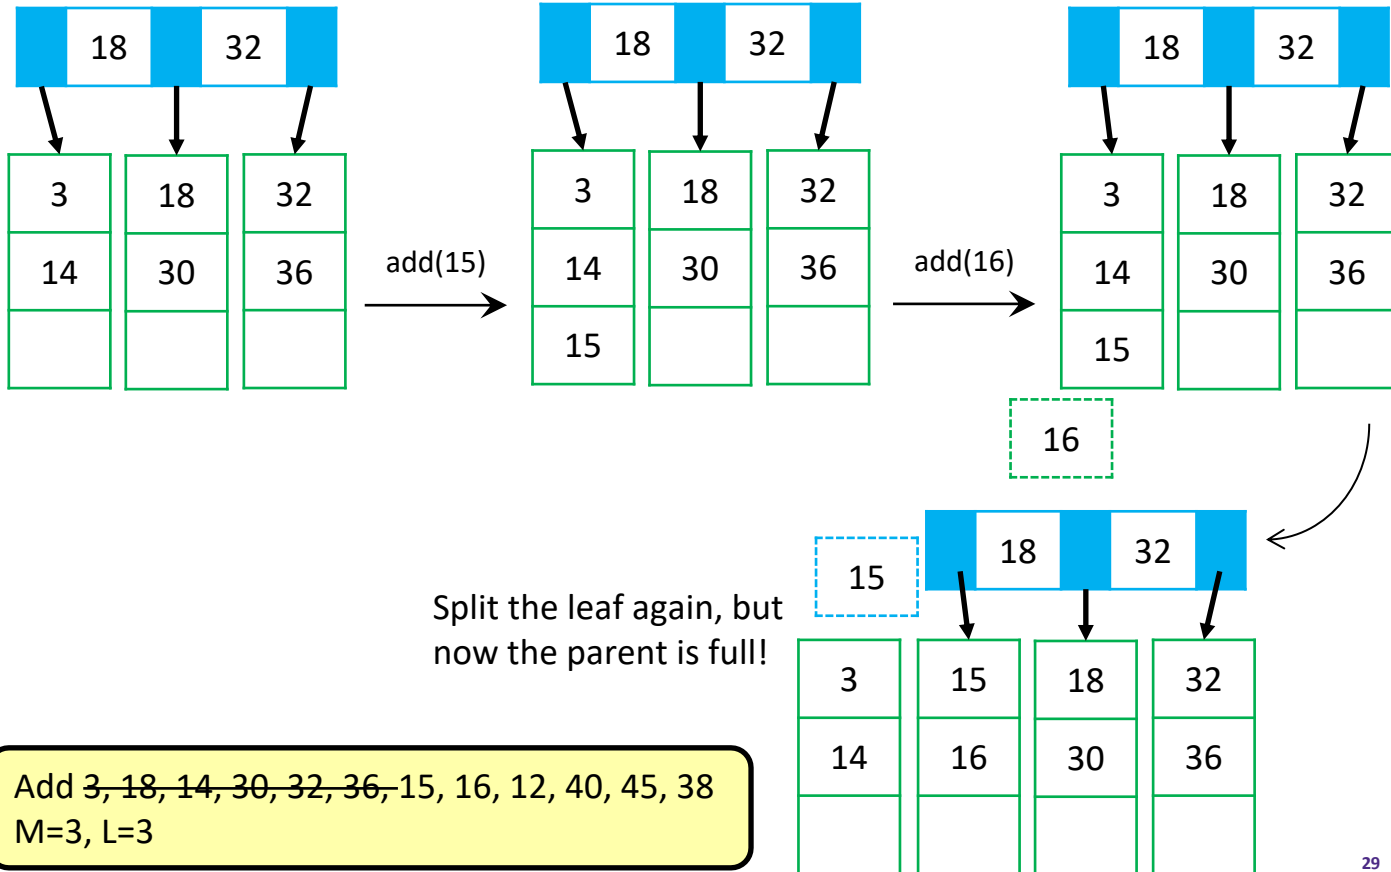
Add ~~3, 18, 14, 30~~, 32, 36, 15, 16, 12, 40, 45, 38
M=3, L=3

Add Example: Answer (3 of 7)




Add ~~3, 18, 14, 30, 32, 36, 15, 16, 12, 40, 45, 38~~
M=3, L=3

Add Example: Answer (4 of 7)




Add Example: Answer (5 of 7)



Split the parent (in this case, the root).
Note that the median key **moves** into the parent (vs being copied)


Add ~~3, 18, 14, 30, 32, 36, 15, 16, 12, 40, 45, 38~~
M=3, L=3

Add Example: Answer (6 of 7)



Add ~~3, 18, 14, 30, 32, 36, 15, 16, 12, 40, 45, 38~~
 M=3, L=3

Add Example: Answer (7 of 7)



Add ~~3, 18, 14, 30, 32, 36, 15, 16, 12, 40, 45, 38~~
M=3, L=3

B+ Tree Add Algorithm (3 of 3)

❖ Note the similarities between the overflow steps:

Split the **leaf** into two leaves:

- Original **leaf** with $\lceil (L+1)/2 \rceil$ smaller items
- New **leaf** with $\lfloor (L+1)/2 \rfloor = \lceil L/2 \rceil$ larger items

Attach the new **leaf** to its parent

- Add a new key (smallest key in new **leaf**) to the parent in sorted order

Split the **internal node** into two leaves:

- Original **node** with $\lceil (M+1)/2 \rceil$ smaller items
- New **node** with $\lfloor (M+1)/2 \rfloor = \lceil M/2 \rceil$ larger items

Attach the new **internal node** to its parent

- Move the median key (smallest key in new **node**) to the parent in sorted order

❖ But also the difference when overflowing a root:

Split the **root** into two **internal nodes**:

- Left **node** with $\lceil (M+1)/2 \rceil$ smaller items
- Right **node** with $\lfloor (M+1)/2 \rfloor = \lceil M/2 \rceil$ larger items

Attach the **internal nodes** to the new **root**

- Move the median key (smallest key in new right **node**) to the **root**

gradescope.com/courses/275833

- ❖ When splitting nodes in a B+ Tree, why do we need to *copy* keys out of leaves but *move* keys out of internal nodes?

B+ Tree Add: Efficiency (1 of 2)

- ❖ Find correct **leaf**: $O(\log_2 M \log_M n)$
- ❖ Add (key, value) pair to **leaf**: $O(L)$
 - Why?
- ❖ Possibly split **leaf**: $O(L)$
 - Why?
- ❖ Possibly split parents all the way up to **root**: $O(M \log_M n)$
 - Why?

- ❖ Total: $O(L + M \log_M n)$

B+ Tree Add: Efficiency (2 of 2)

- ❖ Worst-case runtime is $O(L + M \log_M n)$!
- ❖ But the worst-case isn't that common!
 - Splits are uncommon
 - Only required when a node is full
 - M and L are likely to be large and, after a split, nodes will be half empty
 - Splitting the **root** is extremely rare
 - Remember that our goal is minimizing disk accesses! Disk accesses are still bound by $O(\log_M n)$

Lecture Outline

❖ Recap

❖ B+ Trees

- Goals and Design
- B+ Tree Structure
- B+ Tree Find
- B+ Tree Add
- **B+ Tree Remove**

B+ Tree Remove Algorithm (1 of 3)

1. Remove the item from its **leaf**
2. If the **leaf** now has $\lceil L/2 \rceil - 1$, *underflow*:
 - If a neighbor has $> \lceil L/2 \rceil$ items, *adopt*
 - Move parent's key down, and neighbor's adjacent key up
 - Else, *merge leaf* with neighbor
 - Guaranteed to have a legal number of items
 - Remove parent's key and move grandparent's key down
 - Parent now has one less **leaf**

If step (2) caused the parent to have $\lceil M/2 \rceil - 1$ children, ...

B+ Tree Remove Algorithm (2 of 3)

3. If step (2) caused an **internal node** to have $\lceil M/2 \rceil - 1$ children
 - If a neighbor has $> \lceil M/2 \rceil$ keys, *adopt* and update parent
 - Move parent's key down, and neighbor's adjacent key up
 - Else, *merge* with neighbor node
 - Guaranteed to have a legal number of keys
 - Remove parent's key and move grandparent's key down
 - Parent now has one less node, may need to continue up the tree
4. If step (3) caused the **root** to have have $\lceil M/2 \rceil - 1$ children
 - If **root** went from 2 children to 1 child, move key down and make the child the new **root**
 - *This is the only case that decreases the tree height!*

B+ Tree Remove Algorithm (3 of 3)

❖ Again, note the similarities between the underflow steps:

If a neighbor **leaf** has $> \lceil L/2 \rceil$ items,
adopt:

Move parent's key down, and
neighbor's adjacent key up

Else *merge leaf* with neighbor:

Guaranteed to have a legal
number of items

Remove parent's key and move
grandparent's key down

Parent now has one less **leaf**

If a neighbor **node** has $> \lceil M/2 \rceil$ items,
adopt:

Move parent's key down, and
neighbor's adjacent key up

Else *merge node* with neighbor:


Guaranteed to have a legal number of
keys

Remove parent's key and move
grandparent's key down


Parent now has one less **leaf**

Remove Example

- ❖ Remove 32, 15, 16, 14, 18
- ❖ $M=3, L=3$
 - Min #children = 2
 - Min #items = 2
- ❖ Gradescope question:
 - How many nodes do we end with?




Remove Example: Answer (1 of 8)




remove(32)
→

Remove 32, 15, 16, 14, 18
M=3, L=3; min children=2, min items=2




Remove Example: Answer (2 of 8)




remove(15)
→

Remove ~~32~~, 15, 16, 14, 18
M=3, L=3; min children=2, min items=2




Remove Example: Answer (3 of 8)




remove(16)
→


Remove ~~32, 15, 16~~, 14, 18
M=3, L=3; min children=2, min items=2




Remove Example: Answer (4 of 8)



Remove ~~32, 15, 16~~, 14, 18
 M=3, L=3; min children=2, min items=2




Remove Example: Answer (5 of 8)




remove(14)
→

Remove ~~32, 15, 16, 14, 18~~
M=3, L=3; min children=2, min items=2




Remove Example: Answer (6 of 8)




remove(18)
→


Remove ~~32, 15, 16, 14, 18~~
M=3, L=3; min children=2, min items=2



Remove Example: Answer (7 of 8)




Merge with a neighbor **node**




Remove ~~32, 15, 16, 14, 18~~
M=3, L=3; min children=2, min items=2

Remove Example: Answer (8 of 8)



Delete the old **root**



Remove ~~32, 15, 16, 14, 18~~
 M=3, L=3; min children=2, min items=2

B+ Tree Remove: Efficiency (1 of 2)

- ❖ Find correct **leaf**: $O(\log_2 M \log_M n)$
- ❖ Remove item from **leaf**: $O(L)$
 - Why?
- ❖ Possibly adopt from or merge with neighbor **leaf**: $O(L)$
 - Why?
- ❖ Possibly adopt or merge **parent node** up to **root**: $O(M \log_M n)$
 - Why?
- ❖ Total: $O(L + M \log_M n)$

B+ Tree Remove: Efficiency (2 of 2)

- ❖ Worst-case runtime is $O(L + M \log_M n)$!
- ❖ But the worst-case isn't that common!
 - Merges are uncommon
 - Only required when a node is half empty
 - M and L are likely large and, after a merge, nodes will be completely full
 - Shrinking the height by removing the **root** is extremely rare
 - Remember that our goal is minimizing disk accesses! Disk accesses are still bound by $O(\log_M n)$

Lecture Outline

❖ Recap

❖ B+ Trees

- Goals and Design
- B+ Tree Structure
- B+ Tree Find
- B+ Tree Add
- B+ Tree Remove
- **Wrap-Up**

B+ Trees in Java?

- ❖ For most of our data structures, we encourage writing high-level, reusable code. Eg, using Java generics in our projects
- ❖ It's a bad idea for B+ Trees, however
 - Java can do balanced trees!
 - Java wasn't designed for things like managing disk accesses, which is the whole point of B+ Trees
 - The key issue is Java's extra *levels of indirection*...

Possible Java Implementation: Code



Even if we assume `int` keys, Java's data representation doesn't match what we want out of a B+ Tree

```
class BTreeNode<E> { // internal node
    static final int M = 128;
    int[] keys = new int[M-1];
    BTreeNode<E>[] children = new BTreeNode[M];
    int numChildren = 0;
    ...
}



class BTreeLeaf<E> { // leaf node
    static final int L = 32;
    int[] keys = new int[L-1];
    E[] items = new Object[L];
    int numItems = 0;
    ...
}
```

Why is the code bad for B+Tree?

BTreeNode (internal node)



BTreeLeaf (leaf node)



*All the **red** references indicate “unnecessary” indirection that might be avoided in another programming language!*

B+ Trees in Java: Just say no

- ❖ The whole idea behind B+ trees was to keep related data in contiguous memory
- ❖ But this runs counter to the code and patterns Java encourages
 - Java's implementation of generic, reusable code is not what you want for your performance-critical web-index
- ❖ Other languages (e.g., C++) have better support for “flattening objects into arrays” in a generic, reusable way
- ❖ Levels of indirection matter!

Summary: Search Trees

- ❖ **Binary Search Trees** make good dictionaries because they implement **find**, **add**, and **remove** as well as a number of useful operations such as ~~flatten~~**IntoSortedList** or **successor**
 - Essential and beautiful computer science
- ❖ *Balanced* search trees guarantee logarithmic-time operations
 - ... if you can maintain balance within the time bound
 - **AVL trees** maintain balance by tracking height and allowing all children to differ in height by at most 1
 - **B trees** maintain balance by keeping nodes at least half full and all leaves at same height
- ❖ Next up: dictionaries that don't rely on trees at all!